

Project 3

Deadline:	Submit by midnight of 24 May 2020.
Evaluation:	25% of your final course grade.
Late Submission:	See Course Guide.
Work	This assignment may be done in pairs . No more than two people per group are allowed. Should you choose to work in pairs, upon submission of your assignment, you will need to fill out and submit a form (to be provided) indicating your contribution to the project.
Purpose:	Learning outcomes 1 - 5 from the course outline.

Project outline:

Kaggle (<https://www.kaggle.com/>) is a crowdsourcing, online platform for machine learning competitions, where companies and researchers submit problems and datasets, and the machine learning community compete to produce the best solutions. This is a perfect training ground for real-world problems. It is an opportunity for data scientists to develop their portfolio which they can advertise to their prospective employers, and it is also an opportunity to win prizes.

For this project, you are going to work on a Kaggle dataset.

You will first need to create an account with Kaggle. Then familiarise yourself with the Kaggle platform.

Your task will be to work on a competition dataset which is currently in progress. The problem description and the dataset can be found here <https://www.kaggle.com/c/m5-forecasting-uncertainty/overview>

Task:

Your work is to be done using the Jupyter Notebook, which you will submit as the primary component of your work.

Your tasks are as follows:

1. You will first need to create an account with Kaggle.
2. Then familiarise yourself with the Kaggle platform.
3. Familiarise yourself with the submission process.
4. Download the datasets, then explore and perform thorough EDA.
5. Devise an experimental plan for how you intend to empirically arrive at the most accurate solution.
6. Explore the accuracy of kNN for solving the problem.
7. Explore scikit-learn (or other libraries) and employ a suite of different machine learning algorithms not yet covered in class.
8. Investigate which subsets of features are effective, then build solutions based on this analysis and reasoning.
9. Devise solutions to these machine learning problems that are creative, innovative and effective. Since much of machine learning is trial and error, you are asked to continue refine and incrementally improve your solution. Keep track of all the different strategies you have used, how they have performed, and how your accuracy has improved/deteriorated with different strategies. Provide also your reasoning for trying strategies and approaches. Remember, you can submit up to four solutions to Kaggle per day. Keep track of your performance and consider even graphing them.
10. Take a screenshot of your final and best submission score and standing on the Kaggle leader-board for both competitions and save that as a jpg file. Then embed this jpg screenshots into your Notebooks, and record your submission scores on the class Google Doc (found on Stream) where the class leader-boards will be kept.

The Kaggle platforms and the community of data scientists provide considerable help in the form of 'kernels', which are often Python Notebooks and can help you with getting started. There are also discussion fora which can offer help and ideas on how to go about in solving problems. **Copying code from this resource is not acceptable for this assignment.** Doing so can be regarded as plagiarism, and can be followed with disciplinary action.

Marking criteria:

Marks will be awarded for different components of the project using the following rubric:

Component	Marks	Requirements and expectations
EDA	5	Variety of exploratory research and inquiry into different aspects of the dataset, use of broad and appropriate range of visualisations and their effective communication. Thoroughness in data preparation.
Classification modelling using kNN	30	Experimentation with kNN. Considering different values of k and effects of different distance metrics.
Classification modelling using a variety of algorithms	25	It is unlikely that kNN will produce the best (or even satisfactory) accuracy on this kind of problem. Therefore, you are asked to explore and use a variety of algorithms either from scikit-learn, or elsewhere, in order to arrive at your best solution for the competition.
Analysis	20	The manner in which you have devised your experiments, evaluated your classifiers, interpreted your findings, as well as conducted feature analysis and feature selection.
Kaggle submission score	20	<p>Successful submission of predictions to Kaggle, listing of the score on the class leader-board and position on the class leader-board.</p> <p>The winning student will receive full marks. The next best student will receive 17 marks, and every subsequent placing will receive one less point, with the minimum being 10 marks for a successful submission.</p> <p>An interim solution must be submitted by May 10 and the class leader board document must be updated. This will constitute 10 marks. If this is not completed by this date, then 10 marks will be deducted from the submission score. For this, you must submit a screenshot of your submission date and score.</p>
BONUS MARKS		
Cluster analysis	5	Use of cluster analysis for exploring the dataset.
Additional feature extraction	5	Bonus marks will be awarded for extracting additional features from this dataset and incorporating them into the training set, together with the comparative analysis showing whether or not they have increased predictive accuracy.

Hand-in: Zip-up all your **notebooks**, any other .py files you might have written as well as jpgs of your screenshots into a single file and submit through Stream. **If, and only if** Stream is down, then email the solution to the lecturer.

If you have any questions or concerns about this assignment, please ask the lecturer sooner rather than closer to the submission deadline.